

Annotation Standardization Guide

Dialectic Preference Bias in LLMs Study · Phase 2 · April 2026 · York University / Vector Institute — Seyyed-Kalantari Lab

This document establishes the shared annotation standard agreed upon by the three annotators (A1 — University of Toronto, A2 — University of Guelph, A3 — Wilfrid Laurier University) following Phase 1 disagreement analysis. It covers the two core fields where inter-annotator agreement was lowest: **Sentiment Classification** and **Semantic Equivalence (Meaning Changed)**. All annotators must apply these rules uniformly in Phase 2 across the 400-tweet dataset.

Part 1 — Sentiment Classification

Sentiment must be assessed from the perspective of an **outside observer** reading the tweet cold — not from the speaker's intent. The three valid labels are **Positive**, **Negative**, and **Neutral**. The decision process below must be followed in order.

Core Principles

Observer, not speaker	Enthusiasm ≠ Positive	Ending sentiment wins
Judge how a neutral reader would feel reading the tweet, not how the author felt writing it. "I like murder" is Negative to an observer even if the speaker was positive.	High energy, excitement, or strong language does not automatically mean Positive. AAE often uses expressive register for neutral or negative content.	If the tweet shifts tone midway, the <i>final emotional note</i> determines the label. A tweet that vents then resolves positively = Positive.

Step-by-Step Decision Flow

Work through these questions **in order**. Stop at the first question that gives you a definitive answer.

Ask yourself	YES →	NO →
Does the tweet describe real physical harm, a direct attack on a person, or clearly toxic/harmful intent?	NEGATIVE	Continue to next question
Does the tweet describe toxic behaviour, cruelty, or bad behaviour directed at others?	NEGATIVE	Continue to next question
Is the tweet a plain statement of fact, a logistical update, or a matter-of-fact observation with no emotional landing?	NEUTRAL — see Neutral rules below	Continue to next question

Does the tweet end on a clearly positive note — warmth, celebration, pride, affection, or positive resolution?	POSITIVE	Continue to next question
Does the tweet end on a clearly negative note — frustration, disappointment, complaint, or unresolved tension?	NEGATIVE	Continue to next question
Does the tweet contain positive internet slang (e.g. "slay", "lowkey fire", "no cap it's good") used sincerely?	POSITIVE — unless clearly sarcastic	NEUTRAL — default if no clear valence

Label Definitions with Examples

POSITIVE

The overall emotional landing is good — includes warmth, affection, pride, celebration, gratitude, humor that resolves positively, or statements of positive belief/value.

Positive if:

- The tweet ends on warmth or affection ("he ain't wanna leave his auntie")
- The speaker expresses wanting something good ("I want and need you")
- The tweet contains a positive belief or value statement ("It takes a real man to realize one woman is enough")
- Positive internet slang used sincerely
- Media/lyric references that are celebratory or self-affirming

NEGATIVE

The overall emotional landing is harmful, distressing, or socially damaging — includes real harm, directed cruelty, toxic beliefs, or unresolved frustration/disappointment.

Negative if:

- The tweet describes or endorses real physical harm
- The tweet is directed cruelty or mockery toward a specific person/group
- A toxic word is being used as harm, not as emphasis
- The tweet expresses a toxic belief ("Is this a mean or harmful worldview?")
- The speaker is clearly distressed or disappointed with no resolution

NEUTRAL

The tweet is a matter-of-fact statement, logistical update, objective observation, or question seeking factual information — with no clear emotional landing. **Neutral is not the default for uncertainty** — only use it when the tweet genuinely carries no positive or negative charge.

Neutral if:

- Plain logistical statement: "Im going to sleep wen I get off the phone"
- Objective observation with no emotional stake: "I ain't got cable either. Reason why I'm watching the news"
- Casual invitation with no expressed feeling: "Goin to da movies who tryin to come"
- Factual question: "You going? Or nahh?"
- Two opposing sentiments that cancel each other out with no resolution

NOT neutral just because: the tweet is short, uses flat language, or the subject is ordinary.

Special Cases in AAE

Pattern	How to handle	Example
Sarcasm / Bait-and-switch	Follow the ending sentiment. If it ends negatively after a positive setup, call Negative.	<i>"lol o yea?" after a positive plan = likely Positive</i>
Gender-dependent sentiment	If sentiment depends on the speaker's gender (e.g. "Deets on the D"), choose the reading that is most likely given contextual cues. If truly ambiguous, flag it.	<i>A tweet about physical attraction may read differently for different annotators</i>
Missing cultural context	If you genuinely cannot determine sentiment without cultural context you don't have, flag it rather than guessing. Do not use SAE translation to decode the original AAE.	<i>Sports references, media quotes, or local slang you don't recognize</i>
Media/lyric references	These are rarely neutral. Apply positive or negative based on the content of the reference itself if known, or the framing around it if not.	<i>"Me I'm no mobster, me I'm just me!" = Positive (self-affirmation)</i>
"chill" and similar AAE relaxation terms	"Chill" indicates a mild positive feeling — contentment. It is not neutral unless surrounded by entirely flat context.	<i>"I think I'm just gone chill" = leans Positive</i>

Part 2 — Semantic Equivalence (Meaning Changed)

The question is simple: did the SAE translation convey the **same meaning** as the original AAE, or did it convey a **different meaning**? This study is designed to detect extremes, not gray areas. The standard is **literal one-to-one**: if a native AAE speaker would read the SAE version and say "that's what I said," the meaning is the same. If they would say "that's not what I meant," the meaning has changed.

Use "Partly Same" sparingly — only when the propositional content is preserved but a critical element (e.g. a specific cultural term, an idiomatic phrase with no SAE equivalent) has been approximated rather than translated.

Step-by-Step Decision Flow

Ask yourself	YES →	NO →
Does the original AAE make sense as written? (Is it grammatically interpretable as AAE?)	Continue to next question	MEANING CHANGED — if the original is incoherent, the SAE cannot be a faithful translation
Is the tweet a reply to another tweet (not a standalone statement)?	Note the missing context. If the SAE translation invented a context the AAE didn't have, MEANING CHANGED	Continue to next question
Does the SAE version refer to the same person, object, or event as the original AAE?	Continue to next question	MEANING CHANGED — subject substitution
Do any AAE-specific terms (insults, slang, idioms) carry a cultural weight that is lost or replaced in the SAE?	MEANING CHANGED — see insult/slang rules below	Continue to next question
Would a native AAE speaker read both versions and say they mean the same thing?	SAME MEANING	MEANING CHANGED

Label Definitions with Examples

SAME MEANING

The SAE version conveys the same propositional content and pragmatic intent as the original AAE. Minor spelling corrections, punctuation normalization, and register-neutral word choice differences do not constitute meaning change.

Still same meaning even if:

- Abbreviations are expanded ("philly" → "Philadelphia", "SUCS" → "Buccaneers")
- Spelling is corrected
- Near-synonyms used where cultural weight is equivalent ("lame" → "loser" — *depends on context, see insult rule*)
- Non-AAE informal terms corrected (e.g. "bout me" where "bout" is not functioning as AAE)

MEANING CHANGED

The SAE version says something factually, pragmatically, or culturally different from the original AAE. The test: would a native AAE speaker read both and say "that's not what I meant"?

Meaning has changed when:

- **Subject/object substitution:** original refers to personality, SAE refers to a physical attribute
- **Insult weight lost:** "bitch ass" → "coward" — different social register and community meaning
- **Insult weight lost:** "lame" (AAE: insecure, lacking self-assurance) → "loser" (failing in all aspects) — crucially different
- **AAE idiom flattened:** "Oh my ass!" → "Oh, come on!" — completely different exclamation, different emotional register
- **AAE-specific opener deleted:** "shid nigha" deleted entirely rather than translated
- **Incoherent AAE falsely resolved:** original didn't make sense, SAE invented a meaning
- **Vague AAE term over-specified:** "hoes" → "people/folks" — far too general
- **Idiomatic phrases not translated:** "for a minute" (AAE: for a long time) → literally

The Insult and Slang Weight Rule

AAE insults and slang carry community-specific cultural weight that standard SAE equivalents often miss. Before accepting a near-synonym as "same meaning," ask: *in AAE culture, does this word target the same quality in a person as the SAE equivalent?* If the answer is no, the meaning has changed.

AAE Term	SAE Translation	Ruling	Why
<i>lame</i>	<i>loser</i>	CHANGED	In AAE, "lame" = insecure, lacking self-assurance. "Loser" = failing at life. Different quality being targeted.
<i>bitch ass</i>	<i>coward</i>	CHANGED	"Bitch ass" in AAE carries a specific social register and directness that "coward" does not replicate.
<i>hoes</i>	<i>people / folks</i>	CHANGED	"Hoes" is a specific reference; "people" is too general and loses the speaker's specific framing.
<i>nigga / MF / they</i>	<i>used as self-reference</i>	SAME	In AAE, these terms are often 3rd-person self-references. If the SAE preserves the self-reference meaning, it's same meaning.

Part 3 — Survey Instrument Improvement Recommendations

The following issues were identified by annotators during Phase 1 and the standardization session. These are submitted to the research team (Ryan Wu / Dr. Seyyed-Kalantari) as recommended changes to the annotation instrument before the 400-tweet Phase 2 deployment.

01 Mental Preparation — Cultural Context Flag

Some tweets reference specific cultural events, media, or community figures that may not be familiar to all annotators. Without this context, annotators risk misidentifying meaning change or sentiment when the real issue is missing cultural knowledge. Recommendation: add a pre-annotation flag that allows annotators to mark "I do not have the cultural context for this tweet" before answering, so the research team can weight those items accordingly.

02 Free-Text Explanation After Flagging

The current instrument captures what was flagged but not why. When an annotator flags meaning changed or unnatural back-translation, there is no field to explain the specific feature that triggered the flag. Recommendation: add a short optional text field (50–150 characters) after any non-default answer, allowing annotators to name the specific issue.

03 Blinded Back-Translation Evaluation

Annotators currently evaluate back-translation naturalness while viewing the original AAE, creating an anchoring bias — the original shapes the judgment of the back-translation. Recommendation: provide a toggle to hide the original AAE and SAE when evaluating back-translation naturalness, forcing an independent evaluation on the back-translation's own terms.

04 Bait-and-Switch / Juxtaposed Sentiment Field

AAE frequently uses a bait-and-switch structure where sentiment starts positive and ends negative (or vice versa). The current 3-label scale cannot capture this. Recommendation: add a secondary toggle for "juxtaposed sentiment" that annotators can activate when they detect opposing sentiments within the same tweet, with a short description of what they observed.

05 Back-Translation Unnaturalness Subcategory

Currently, "unnatural" covers two very different failure types: (a) logically incoherent AAE, and (b) AAE markers misplaced or overused — the "AI performing AAE" problem. These have different implications for the study. Recommendation: split the unnatural option into "incoherent AAE" and "misplaced/overused AAE markers" as separate options.

06 "Cannot Confirm" Option for Meaning

The current instrument forces a binary or ternary decision on meaning even when the annotator genuinely cannot tell. Recommendation: add a "Cannot confirm — context missing" button that feeds into a separate adjudication queue rather than being treated as a valid annotation label.

0
7

Explicit Gender Perspective Prompt

Sentiment can vary meaningfully depending on the inferred gender of the speaker or addressee in the tweet. The current instrument provides no prompt for this. Recommendation: add an optional field asking annotators to note if gender context affected their sentiment judgment and how.

0
8

Annotator Familiarity Signal Per Item

Background familiarity is currently captured once at signup, but it varies item by item. Recommendation: add a quick 3-point per-item familiarity slider (unfamiliar / somewhat familiar / fully familiar) so the research team can weight annotations by item-level familiarity rather than only by annotator-level background.